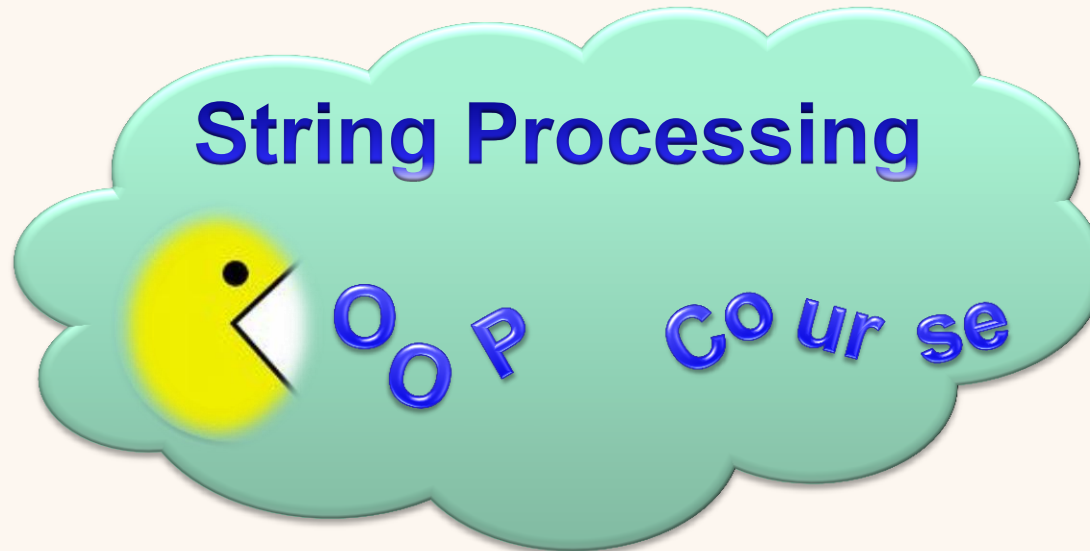


Introduction to Object Oriented Programming

(Hebrew University, CS 67125 / Spring 2014)

Lecture 11



Question

```
/**
```

```
* A method that return true iff the given string is a legal  
* date of the format dd/MM/yyyy
```

```
*/
```

```
boolean isDate(String s) {  
    ???  
}
```


- 01/05/2012
- 15/15/1214
- Hello
- 1a/02/2011
- 111/01/2012

Regular Expressions

- A regular expression is a kind of **pattern** that can be applied to text (*Strings*, in Java)
- A regular expression either **matches the text** (or part of the text), or it **fails** to match it
 - **Which part** of the text matches?
 - **Which parts** of the regular expression match **which parts** of the matched text?
 - Can we **perform substitutions** on the text?
- Regular expressions are an extremely useful tool for manipulating text

Regular Expressions Uses

- Text search
- Filtering/mining/tagging text data (books/forums/blogs)
- Automatic generation of Web pages
- Processing database queries
- Spelling
- Summarization



-

How Do They Look Like?

`[A-Za-z]+ (or)? .* $`

`^[^A]+a+[^a]+`

`(.*) .+ (.*)`

`^[ABC]+\d*$`

`Cats? and rats?`

`(\d+)_ (\d+). (\d+)`

`[a-z]{1,4}?[A-Z1-9]`

Basic Syntax Characters

Char	Usage	Example
a,b,c,...	Regular text	abc matches abc
.	Matches any single character	.at matches cat, bat, rat, 1at...
[...]	Matches any single character of the ones contained	[cbr]at matches cat, bat, rat.
[^...]	Matches any single character except for the ones contained	- [^bc]at matches rat, sat... , but does not match bat, cat.
[a-z]	Matches any character in the range a-z Also works for A-Z, 0-9, and in the negative form (with ^)	- [f-l]aaa matches faaa, gaaa,...,laaa - [^a-f]aaa matches gaaa, 5aaa, &aaa , but does not match aaaa, caaa,
...		...

Basic Syntax

Quantifiers

Char	Usage	Example
*	Matches zero or more occurrences of the single preceding character	<ul style="list-style-type: none">- .<i>*at</i> matches everything that ends with <i>at</i>: <i>at</i>, <i>hat</i>, <i>123_\$&treat...</i>- <[<i>^</i>>]*> matches <i><...anything...></i>
+	Matches one or more occurrences of the single preceding character	0+123 matches <i>0123</i> , <i>00123</i> , <i>000123...</i>

How to catch a rabbit?



Goal: detect in the text all the instances of the word “rabbit”, where each letter may be duplicated many times

Example: rrrrrrabbit, rrabbit, rabbbit, rraabbiitt, rabbbit

Regexp:

`r+a+bb+i+t+`

Search Types

- Consider the regular expression "[a-z]+"
- This expression will match a sequence of one or more lowercase letters
 - [a-z] means any character from a through z, inclusive
 - + means **“one or more”**
- Suppose we apply it on the String "Now is the time"

Search Types

"[a-z]+" , "Now is the time"

- There are *two ways* we can apply this pattern:
 - Apply it on the ***entire string***:
 - **Failure:** The string contains characters other than **lowercase letters**
 - To search the string ***sequentially*** (i.e. search for a substring which matches the given pattern):
 - **Success:** match **ow**
 - If applied repeatedly, it will find **is**, then **the**, then **time**, then **fail**

How to catch a rabbit?



Goal: detect in the text all the instances of words ending with “rabbit”, where each letter except the last may be capitalized

Example: MyRabbit, HISRABBIT, RabBit, gOOdrABBIT

Regexp:

[A-Za-z]*[Rr][Aa][Bb][Bb][lI]t

Recursive Structure

- The structure of regular expressions is **recursive**
 - The vertical bar, `|`, is used to separate alternatives
 - For example, the pattern `abc|xyz` will match either `abc` or `xyz`
 - If one pattern is followed by another, the two patterns must match consecutively
 - For example, `[A-Za-z]+[0-9]` will match one or more letters (`[A-Za-z]+`) immediately followed by one digit (`[0-9]`)
 - Both these operators can be used recursively
- Parentheses may be used in building the expressions
 - For example, `(a|b)c` will search either `a` or `b`, followed by `c`

Predefined Character Classes

.	Any character (may or may not match line terminators)
\d	A digit: [0-9]
\D	A non-digit: [^0-9]
\s	A whitespace character: [\t\n\x0B\f\r]
\S	A non-whitespace character: [^\s]
\w	A word character: [a-zA-Z_0-9]
\W	A non-word character: [^\w]

Notice the space.
Spaces are significant
in regular expressions!

How to catch a rabbit?



Goal: detect in the text all the instances of “1/one rabbit” or “X rabbits” where X is a number
Example: one rabbit, 256 rabbits

Regexp:

(1|one) rabbit|\d+ rabbits

Boundary Matchers

<code>^</code>	The beginning of a line
<code>\$</code>	The end of a line
<code>\b</code>	A word boundary
<code>\B</code>	A non-word boundary

Example:

REGEX is: `\brabbit\b`

INPUT is: Some rabbits plays in the yard.

No match found.



Quantifiers

- Assume X represents some pattern
 - X may be a **single character**, a **range of characters** or an **expression wrapped by parentheses**
 - $X\{n\}$ X occurs exactly n times
 - $X\{n,\}$ X occurs n or more times
 - $X\{n,m\}$ X occurs at least n but not more than m times
 - $X?$ X is optional (it occurs once or not at all) [$\Leftrightarrow X\{0,1\}$]
 - X^* X occurs zero or more times [$\Leftrightarrow X\{0,\}$]
 - X^+ X occurs one or more times [$\Leftrightarrow X\{1,\}$]
- These are *postfix* operators (comeing **after** the operand)

How to catch a rabbit?



Goal: detect in the text all the instances of the word rabbit, where each letter is duplicated at most one time and there is a possible single “s” at the end

Example: rrabbbbit, rabbbits

Regexp:

$r\{1,2\}a\{1,2\}b\{2,4\}i\{1,2\}t\{1,2\}s?$

Types of Quantifiers

- A **greedy quantifier** will match **as much as it can**, and **back off** if it needs to

$X?$ X^* X^+ $X\{n\}$ $X\{n,\}$ $X\{n,m\}$

- A **reluctant quantifier** will match **as little as possible**, then take more if it needs to

- You make a quantifier reluctant by appending a '?':

$X??$ $X^*?$ $X^+?$ $X\{n\}?$ $X\{n,\}?$ $X\{n,m\}?$

- A **possessive quantifier** will match as much as it can, and **never let go**

- You make a quantifier possessive by appending a '+':

$X?+$ X^*+ X^++ $X\{n\}+$ $X\{n,\}+$ $X\{n,m\}+$




Examples

Greedy Quantifier

will match as much as it can, and back off if it needs to

- Current INPUT is: xfooxxxxxxfoo

– Current REGEX is: .*foo

- Searching for .*
- Searching for foo 
- Going back one letter
- Searching for foo 
- Going back another letter
-
- Found foo! 

xfooxxxxxxfoo



- **"xfooxxxxxxfoo"** found starting at index [0,13]

Examples

Reluctant Quantifier

*will match **as little as possible**, then take more if it needs to*

- Current INPUT is: xfooxxxxxfoo

– Current REGEX is: .***?**foo

- Searching for .*
- Searching for foo 😞
- Grabbing another character
- Searching for foo 😊
- Searching for .* again
- Searching for foo again 😞
- Grabbing another character
-
- Found foo! 😊

xfooxxxxxfoo



- "xfoo" found starting at index [0,4]
- "xxxxxxfoo" found starting at index [4,13]

Examples

Possessive Quantifier

*will match as
much as it can,
and **never let go***

- Current INPUT is: xfooxxxxxxfoo

- Current REGEX is: .***+**foo

- Searching for .*
 - Searching for foo



xfooxxxxxxfoo



- No match found

Backreferences

- Say we want to look for a letter that appears twice in a row
 - aa, bb, CC, ...
- How about [a-zA-Z]{2}?
 - This pattern looks for two letters, **not necessarily the same**
 - ab, BC, Aa, ...
- To look for the **same letter** twice, we need to use capturing groups

Capturing Groups

- Parentheses **capture** anything matched by that part of the pattern
 - Example: `([a-zA-Z]*)([0-9]*)`
 - If match succeeds, `\1` holds the matched **letters** and `\2` holds the matched **digits**
 - We can access these parts later in the expression
- `([a-zA-Z])\1` will match a double letter, such as letter

How to catch a rabbit?



Goal: Detect a 'rabbit' with the same number before and after him

Example: 1rabbit1,123rabbit123

Regexp:

(\d+)rabbit\1

Capturing Groups

- Capturing groups are numbered by **counting their *opening parentheses*** from left to right:

(₁ (₂ A) (₃ B (₄ C)))

\1 = ((A)(B(C))), \2 = (A), \3 = (B(C)), \4 = (C)

Spaces

- There is only one thing to be said about white spaces (blanks) in regular expressions, but it's important:

Spaces are significant!

- A space stands for a *space* – when you put a space in a pattern, it means you want to match a **space** in the text string
- It's a *really bad idea* to put spaces in an expression just to make it look better

More Than One Way To Do It

- There is more than one way to write (almost) all regular expressions
 - Especially complex ones
- This applies to mere syntax changes
 - $[a-c]^+ \Leftrightarrow [abc]^+ \Leftrightarrow [abc][abc]^* \Leftrightarrow \dots$
- But also to equivalence between **different** expressions
 - $[ab]\{2\} \Leftrightarrow aa|ab|ba|bb$

Writing Good Regexp

- Writing simple expressions
 - `[ab]{1,2}` is better than `a|b|aa|ab|ba|bb`
- Search for the minimal expression you need
 - If you are only looking for numbers, use `[0-9]*` (or `\d*`), not `.*`
- Anchors are good!
 - When looking for text at the beginning/end of the line, use `^/$`
- Use possessive quantifiers whenever it is possible
- More in

<http://www.tideway.com/confluence//display/Configipedia/Writing+Efficient+Regex>

Thinking in Regular Expressions

- Regular expressions are *not* easy to use at first
 - It's a bunch of punctuation, not words
 - The individual pieces are not hard, but it takes practice to learn to put them together correctly
- They form a miniature programming language
 - It's a different kind of programming language than Java, and requires you to learn new thought patterns
- Despite all this, regular expressions bring much a lot of **power** and **convenience** to String manipulation

Doing It in Java – Preparation

- First, you must ***compile*** the pattern

```
import java.util.regex.*;  
Pattern patt = Pattern.compile("[a-z]+");
```

- Next, you must create a ***matcher*** for a specific piece of text by sending a message to your pattern

```
Matcher matcher = patt.matcher("Now is the time");
```



Pattern and Matcher

- Comments:
 - **Pattern** and **Matcher** are both in `java.util.regex`
 - Neither **Pattern** nor **Matcher** has a public constructor; you create these by using methods in the **Pattern** class
 - A **Matcher** object contains information about *both* the **pattern** to use *and* the **text** to which it will be applied
 - One compiled **Pattern** can be used for many **Matchers**



Doing It in Java – Matching

- Now that we have a Pattern *p* and a Matcher *m*
 - *m.matches()* returns **true** \Leftrightarrow *p* matches the **entire** text string
 - *m.find()* returns **true** \Leftrightarrow *p* matches **any part** of the text string
 - If **called again**, *m.find()* will start searching from where the last match was found
 - *m.find()* returns **true** for as many matches as there are in the string; after that, it returns **false**



Finding What Was Matched

- After a successful match
 - `m.start()` will return the index of the first character matched
 - `m.end()` will return the index of the last character matched, *plus one*
 - These values correspond to what most String methods require
 - For example, `str.substring(matcher.start(), matcher.end())` returns the exact matched substring
- If no match was attempted, or if the match was unsuccessful, `m.start()` and `m.end()` will throw an `IllegalStateException`
 - This is a `RuntimeException`, so you don't have to catch it

A Complete Example

```
import java.util.regex.*;

public class RegexTest {
    public static void main(String args[]) {
        String patternString = "[a-z]+";
        String text = "Now is the time";
        Pattern pattern = Pattern.compile(patternString);
        Matcher matcher = pattern.matcher(text);
        while (matcher.find()) {
            System.out.print(text.substring
                             (matcher.start(), matcher.end()) + "_");
        }
    }
}
```

Output: ow_is_the_time_

Java Shortcuts for Regexp for String Class

<u>Return value</u>	<u>Method name and description</u>
boolean	<i>matches</i> (String regex) Tells whether or not this string matches the given regular expression
String []	<i>split</i> (String regex) splits this string around matches of the given regular expression
String	<i>replaceAll</i> (String regex, String replacement) Replaces each substring of this string that matches the given regular expression with the given replacement

+ Convenient

- Inefficient & inflexible

Question

```
/**
```

```
* A method that return true iff the given string is a legal  
* date of the format dd/MM/yyyy
```

```
*/
```

```
boolean isDate(String s) {
```

```
    ???
```

```
}
```

Question

```
/**
```

```
* A method that return true iff the given string is a legal  
* date of the format dd/MM/yyyy
```

```
*/
```

```
boolean isDate(String s) {  
    return s.matches("\\d{2}/\\d{2}/\\d{4}");  
}
```



So Far...



- **Regular expressions** are a very powerful tool to analyze and manipulate text
 - Basic expressions: `.`, `*`, `+`, `[a-z]`, `a{n,m}`, `w?`, ...
 - Recursive: `expr`, `expr1|expr2`, `expr1expr2`
 - Capturing `(exp1)((exp2)|(exp3))\1\2`
- Patterns in java

```
Pattern patt = Pattern.compile("[a-z]+");
Matcher matcher = patt.matcher("Now is the time");
matcher.matches(), matcher.find()
```